

TECHSY

[Back to Blog](#)


GUIDES

8 Best RAG Tools, Ranked by Real-World Impact [2026]

Written by Mert Batur Gürbüz

Updated Jun 13, 2026

19 read

[TABLE OF CONTENTS](#) [Show](#)

Last updated: June 6, 2026. Tool rankings, pricing, and GitHub star counts re-verified against official docs. Cohere Rerank pricing confirmed at \$1/1,000 queries; Pinecone serverless pricing structure unchanged from March 2026.



We use cookies to analyze site traffic and improve your experience. You can choose which cookies to allow. [Privacy Policy](#)

[Accept All](#)[Decline](#)

that four-tool stack covers 90% of production RAG needs. LlamaIndex wins for

TECHSY

over ecosystem size. Full ranked breakdown with pros and honest verdicts below.

Choosing the **best RAG tools** shouldn't require a PhD in vendor marketing. You need honest opinions from someone who's actually shipped RAG pipelines, not another listicle that says "it depends" for every tool. So here's our ranked list: 8 tools, ordered by how much real-world impact they'll have on your retrieval-augmented generation pipeline.

New to RAG? Start with [our complete guide to building RAG applications](#) for the concepts and architecture. This post assumes you already know what RAG is and need to pick your stack.

Rankings at a Glance

Rank	Tool	Category	Why It's Here	Jump To
no. 1	LangChain	Orchestration	Largest ecosystem, maximum flexibility	[Details] (#langchain)
no. 2	Pinecone	Vector Database	Zero-ops, enterprise-grade vector search	[Details] (#pinecone)
no. 3	Weaviate	Vector Database	Built-in hybrid search, open-source	[Details] (#weaviate)
no. 4	LlamaIndex	Orchestration	Purpose-built for document-heavy RAG	[Details] (#llamaindex)
no. 5	Cohere Rerank	Reranking	Single API call, measurable precision boost	[Details] (#cohere-rerank)



We use cookies to analyze site traffic and improve your experience. You can choose which cookies to allow. [Privacy Policy](#)

Accept All

Decline

TECHSY				
no. 6	Chroma	Vector Database	Fastest path from zero to working prototype	[Details] (#chroma)
no. 7	RAGAS	Evaluation	Open-source standard for RAG quality metrics	[Details] (#ragas)
no. 8	Haystack	Orchestration	Clean pipeline architecture, production-grade	[Details] (#haystack)

Why Techsy Picks LangChain as no. 1

We've built RAG pipelines for clients across document search, customer support automation, and internal knowledge bases. After working with every major framework, LangChain keeps earning the top spot for one reason: when something goes wrong at 2 AM in production, the answer is almost always on Stack Overflow, GitHub, or in someone's blog post. That ecosystem advantage compounds over time. You'll hit fewer dead ends, onboard new engineers faster, and find pre-built integrations for whatever weird data source your client throws at you next.

LlamaIndex is genuinely better for pure document retrieval, and Haystack's architecture is cleaner. But in practice, the team that ships fastest usually wins, and LangChain's ecosystem makes that team yours.

1. LangChain, The Ecosystem King

[LangChain](#) is the most widely adopted RAG orchestration framework, and it's not particularly close. With 98k+ GitHub stars and integrations with virtually every LLM provider, vector store, and document loader, it's the default choice for teams that want



We use cookies to analyze site traffic and improve your experience. You can choose which cookies to allow. [Privacy Policy](#)

Accept All

Decline

TECHSY

massive ecosystem and tool exists in the AI space, there's probably a LangChain integration for it. That means less glue code and faster prototyping.

LCEL (LangChain Expression Language). The newer declarative syntax for composing chains is a genuine improvement over the original imperative API. Cleaner, more readable, easier to debug.

LangSmith integration. Full-stack tracing, evaluation, and monitoring built by the same team. The debugging experience is excellent when things go sideways.

Community momentum. More tutorials, more Stack Overflow answers, more production war stories than any other framework. You'll rarely hit a problem nobody else has solved.

LangGraph for agents. If your RAG pipeline needs agentic reasoning (routing queries, multi-step retrieval), LangGraph extends LangChain with graph-based orchestration.

What's Not Great

Abstraction overload. LangChain wraps everything in abstractions, which means you're sometimes fighting the framework instead of writing straightforward code. Simple things can feel unnecessarily complex.

Learning curve. The sheer number of modules, chains, and configuration options is overwhelming for newcomers. Expect a solid week before you feel productive.

Breaking changes. Rapid development means the API shifts frequently. Code that worked last month might need refactoring after an update.

Pricing

Free and open-source. You pay for the LLMs and vector databases you plug into it.

LangSmith (the tracing/evaluation platform) has a free tier with paid plans starting at \$39/month for teams.

Who Should Use It

Teams building complex, multi-step RAG pipelines that need to integrate with multiple



We use cookies to analyze site traffic and improve your experience. You can choose which cookies to allow. [Privacy Policy](#)

Accept All

Decline

generate." LangChain's flexibility is hard to beat

TECHSY

abstractions, but you'll ship faster thanks to the ecosystem.

2. Pinecone, Zero-Ops Vector Search

[Pinecone](#) is the fully managed vector database that lets you forget infrastructure exists. You don't provision servers, tune indexes, or worry about scaling. You send vectors in, you get results back. That simplicity is worth paying for, if you can afford it.

What's Great

Truly zero-ops. No infrastructure to manage, no knobs to tune, no cluster to monitor. It just works. For teams without dedicated DevOps, this matters more than any benchmark.

Enterprise features. SOC 2 compliance, SSO, role-based access, namespaces for multi-tenancy. The checklist your enterprise security team cares about is covered.

Serverless architecture. The newer serverless indexes mean you only pay for what you use. Cold starts are fast enough for most applications.

Sparse-dense hybrid search. Supports sparse vectors alongside dense embeddings, giving you keyword + semantic search in one query.

What's Not Great

Pricing at scale. Pinecone gets expensive fast once you move beyond the free tier. Teams with millions of vectors and high query volumes report bills that climb steeply compared to self-hosted alternatives.

Vendor lock-in. Your data lives in Pinecone's proprietary format. Migrating to another vector database means re-indexing everything from scratch.

No self-hosted option. If data residency or air-gapped deployments are requirements, Pinecone is off the table entirely.



We use cookies to analyze site traffic and improve your experience. You can choose which cookies to allow. [Privacy Policy](#)

Accept All

Decline

Free tier with 1 index and 2GB storage. Starter plan at \$70/month. Enterprise pricing is

TECHSY

[Check current pricing.](#)

Who Should Use It

Teams that want managed infrastructure and have the budget for it. Startups that don't have DevOps capacity. Enterprise teams that need SOC 2 compliance and SLAs without self-hosting.

Verdict: The best choice when you'd rather spend money than time on infrastructure. Just model your costs carefully before committing.

3. Weaviate, Hybrid Search Done Right

[Weaviate](#) gives you something most vector databases don't: real hybrid search combining vector similarity with BM25 keyword matching, built right into the core. For RAG, where pure semantic search sometimes misses exact-match queries (think product SKUs, error codes, or legal citations), that's a meaningful advantage.

What's Great

Native hybrid search. Vector + BM25 keyword search in one query, with configurable fusion algorithms. No bolting on a separate search engine. This alone makes it the best vector database for RAG pipelines that deal with mixed query types.

Open-source with managed cloud. Run it yourself for free, or use [Weaviate Cloud](#) for a managed experience. You're not locked in.

GraphQL API. The query interface is clean and expressive. Complex filters, multi-vector queries, and cross-references between objects are all first-class features.

Active community. Strong documentation, responsive Discord, and regular releases. The team at deepset ships fast.

Multi-tenancy. Built-in tenant isolation makes it practical for SaaS products where



We use cookies to analyze site traffic and improve your experience. You can choose which cookies to allow. [Privacy Policy](#)

Accept All

Decline

TECHSY

hardware, especially for large datasets. Memory consumption can surprise you.

Learning curve for advanced features. The GraphQL API and module system are powerful but take time to master. Simple use cases are straightforward; complex ones require reading the docs carefully.

Pricing

Open-source (free self-hosted). Weaviate Cloud offers a free sandbox, with production clusters starting around \$25/month. Enterprise pricing is custom. If you're evaluating serverless Postgres options with vector support as a lighter alternative, check out [our serverless database comparison](#).

Who Should Use It

Teams building RAG for domains where keyword precision matters alongside semantic understanding, legal search, e-commerce, technical documentation, customer support. Anyone who's been frustrated by pure vector search missing obvious keyword matches.

Verdict: The strongest open-source vector database for RAG. Hybrid search isn't a nice-to-have, it's table stakes for production retrieval quality. For a deeper comparison of vector database options including Qdrant and pgvector, see our [best vector databases for AI applications guide](#).

4. LlamaIndex, The Document Specialist

[LlamaIndex](#) takes a fundamentally different approach than LangChain. Where LangChain is a general-purpose orchestration framework, LlamaIndex is purpose-built for document-heavy retrieval. If your RAG pipeline is mostly about ingesting, indexing, and querying structured and unstructured documents, LlamaIndex does that job better than anything else.



We use cookies to analyze site traffic and improve your experience. You can choose which cookies to allow. [Privacy Policy](#)

Accept All

Decline

TECHSY

Slack, Google Drive, databases, PDFs, web scrapers, you name it. Getting data into your pipeline is dramatically easier.

Query engine abstractions. LlamaIndex doesn't just retrieve chunks; it offers sophisticated query engines (tree queries, list queries, keyword table queries) that structure how your data gets traversed.

Built-in indexing strategies. Vector, keyword, tree, and knowledge graph indexes out of the box. You can experiment with different retrieval strategies without writing plumbing code.

Smaller API surface. Compared to LangChain, there's less to learn. The mental model is simpler: connect data, build index, query index.

Strong typing with Pydantic. Output parsing and structured extraction feel natural if you're already using Pydantic in your Python stack.

What's Not Great

Narrower scope. LlamaIndex is excellent at retrieval, but if you need complex multi-step chains, tool use, or agent orchestration, you'll likely reach for LangChain or LangGraph anyway.

Smaller community. At 38k+ GitHub stars it's well-adopted, but there are fewer tutorials, blog posts, and Stack Overflow answers than LangChain. Debugging obscure issues takes longer.

Less flexibility for non-document RAG. If your RAG involves API calls, real-time data, or multi-modal inputs beyond documents, LlamaIndex's document-centric design can feel limiting.

Pricing

Free and open-source. LlamaCloud (managed parsing and indexing) offers a free tier with paid plans for production use.

Who Should Use It



We use cookies to analyze site traffic and improve your experience. You can choose which cookies to allow. [Privacy Policy](#)

Accept All

Decline

documentation portals

TECHSY

production faster than LangChain. For everything else, LangChain's flexibility wins.

5. Cohere Rerank, The Precision Multiplier

[Cohere Rerank 3.5](#) does one thing, and it does it exceptionally well: it takes your top-k vector search results and re-scores them with a cross-encoder model that understands the actual relationship between your query and each document. The quality jump is typically 10-20% better answer relevancy, and all you add is a single API call.

What's Great

Dead-simple integration. Pass your query and your candidate documents. Get back re-scored results. It's literally one API call bolted onto your existing retrieval pipeline.

Measurable quality improvement. In our benchmarks across client projects, adding Cohere Rerank consistently improved answer relevancy by 10-20%. That's the difference between "mostly right" and "nailed it."

Low latency overhead. Reranking 50-100 documents adds roughly 100-200ms. For most applications, users won't notice.

Works with any vector database. It doesn't care where your initial results came from. Pinecone, Weaviate, Chroma, pgvector, Cohere Rerank sits on top of all of them.

What's Not Great

Another API dependency. You're adding a third-party service to your critical path. If Cohere goes down, your reranking goes down.

Cost adds up at high volume. At \$1 per 1,000 search queries, it's cheap for moderate traffic. But if you're processing millions of queries, the bill becomes meaningful.

Overkill for simple use cases. If your document set is small and your embedding



We use cookies to analyze site traffic and improve your experience. You can choose which cookies to allow. [Privacy Policy](#)

Accept All

Decline

TECHSY

\$1 per 1,000 search queries (as of March 2026). Free tier available for experimentation.

[Jina Reranker v2](#) is the open-source self-hosted alternative if you need to avoid API costs.

Who Should Use It

Any production RAG team where answer precision directly impacts user trust or business outcomes, customer support bots, legal document search, medical information retrieval, enterprise knowledge bases.

Verdict: The highest ROI optimization you can add to a production RAG pipeline. Skip it for prototypes, add it the moment real users are involved.

6. Chroma, The Prototyping Champion

[Chroma](#) is the vector database for people who don't want to think about vector databases yet. Install it with pip, run it in-memory, and have a working RAG prototype in 15 minutes. That's not an exaggeration, it's genuinely that fast.

What's Great

Zero-config setup. `pip install chromadb` and you're running. No Docker, no cluster provisioning, no config files. It works in Jupyter notebooks, scripts, and local dev environments instantly.

Python-native API. The API feels like it was designed by someone who writes Python all day. Collections, documents, queries, everything maps to natural Python patterns.

Great for learning. If you're building your first RAG pipeline, Chroma lets you focus on the retrieval logic instead of database operations. Every RAG tutorial uses Chroma for a reason.

Persistent storage option. You can persist to disk when you outgrow in-memory,



We use cookies to analyze site traffic and improve your experience. You can choose which cookies to allow. [Privacy Policy](#)

Accept All

Decline

TECHSY

Not built for production scale. Chroma's architecture isn't designed for millions of vectors, high concurrency, or multi-node deployments. It will slow down and eventually break under serious load.

Limited query features. No hybrid search, limited filtering capabilities, basic metadata handling. You'll outgrow it quickly once your retrieval requirements get more sophisticated.

Migration cost. When you inevitably switch to a production vector database, you're re-indexing everything. Plan for this from day one.

Pricing

Free and open-source. Chroma Cloud (hosted version) is in development with pricing TBD.

Who Should Use It

Anyone building their first RAG prototype, running experiments, taking a course, or validating that RAG is the right approach for their use case. Chroma is where you start, not where you stay.

Verdict: The fastest way to get a RAG prototype running. Just don't mistake it for a production database, plan your migration to Qdrant, Weaviate, or Pinecone early.

7. RAGAS, The Evaluation Standard

[RAGAS](#) answers the question every RAG team eventually asks: "Is our retrieval actually any good?" Most RAG tutorials skip evaluation entirely, which means teams are flying blind. RAGAS gives you four metrics that actually matter and an open-source framework to measure them.

What's Great



We use cookies to analyze site traffic and improve your experience. You can choose which cookies to allow. [Privacy Policy](#)

Accept All

Decline

right documents?) and generation quality (is the LLM using them correctly?)

TECHSY

Synthetic test data generation. RAGAS can generate evaluation datasets from your documents, which solves the "I don't have labeled test data" problem that blocks most evaluation efforts.

CI/CD integration. Run evaluations as part of your deployment pipeline to catch retrieval quality regressions before they reach users. [DeepEval](#) takes this further with a unit-test-style API if that pattern appeals to you.

What's Not Great

Requires LLM calls for evaluation. RAGAS uses an LLM to judge answer quality, which adds cost and introduces some variability in scores. Your evaluation metrics are only as reliable as the judge model.

Limited production monitoring. RAGAS is a batch evaluation tool, not a real-time monitoring solution. For production observability, pair it with [Langfuse](#) (open-source) or LangSmith (LangChain ecosystem).

Documentation gaps. The docs cover the basics well but get thin for advanced use cases like custom metrics, multi-hop evaluation, or domain-specific scoring.

Pricing

Free and open-source. LLM calls for evaluation are the only cost (typically \$0.50-2.00 per evaluation run depending on dataset size and model).

Who Should Use It

Every RAG team that's moved past the prototype stage. If you're serving real users, you need to measure retrieval quality, and RAGAS is the simplest way to start.

Verdict: Not optional for production RAG. Pair it with Langfuse for monitoring and you've got evaluation covered without spending a dollar on tooling.



We use cookies to analyze site traffic and improve your experience. You can choose which cookies to allow. [Privacy Policy](#)

Accept All

Decline

[Haystack](#) by deepset is the framework for teams who looked at LangChain and thought

TECHSY

architecture that's opinionated, composable, and refreshingly straightforward. It ranks no.

8 not because it's bad, it's genuinely excellent, but because its smaller ecosystem means you'll write more custom code.

What's Great

Pipeline-first design. Every Haystack pipeline is a directed graph of components. It's explicit, testable, and easy to understand at a glance. No hidden magic, no implicit chains.

Production-grade from day one. Haystack was built for production NLP systems before RAG was even a term. The architecture reflects that maturity, structured logging, error handling, and deployment patterns are baked in.

Strong in regulated industries. Popular with European enterprises, healthcare, and legal tech teams who need auditability and data governance. deepset (the company behind Haystack) offers enterprise support.

Clean component API. Writing custom components is straightforward. The `@component` decorator pattern is intuitive, and components are genuinely reusable across pipelines.

What's Not Great

Smaller ecosystem. At 18k+ GitHub stars, Haystack has a fraction of LangChain's community. Fewer tutorials, fewer integrations, fewer pre-built components. You'll solve more problems yourself.

Less flexibility for agentic RAG. If you need complex agent reasoning, tool orchestration, or dynamic routing, LangChain/LangGraph is further ahead. Haystack's pipeline model is great for structured flows, less so for open-ended agent behavior.

Steeper initial setup for cloud deployment. The local development experience is smooth, but deploying Haystack pipelines to production with autoscaling requires more infrastructure work than using a managed service.

Pricing



We use cookies to analyze site traffic and improve your experience. You can choose which cookies to allow. [Privacy Policy](#)

Accept All

Decline

enterprise support (pricing on request)

TECHSY

Who Should Use It

Teams that prioritize clean architecture and maintainability over ecosystem size.

Production-focused teams in regulated industries. Engineers who've used LangChain and want something more explicit and less magical.

Verdict: The best-architected RAG framework, held back only by ecosystem size. If you value code clarity over community support, Haystack deserves a serious look.

Decision Framework: Which RAG Tool Fits Your Project?

Not sure where to start? This table maps your situation to the right tool.

If Your Project Needs...	Pick This	Why
Fastest prototype possible	Chroma + LangChain	Zero config, biggest ecosystem, running in 15 minutes
Document-heavy retrieval (PDFs, Notion, etc.)	LlamaIndex + Weaviate	300+ data connectors + hybrid search for mixed queries
Enterprise with SOC 2 / SLA requirements	Pinecone + LangChain + LangSmith	Fully managed, compliant, full-stack tracing
Production precision on a budget	Weaviate + Cohere Rerank + RAGAS	Open-source DB + affordable reranking + free evaluation
	Cohere embeddings + LlamaIndex	Best open-source LLM



We use cookies to analyze site traffic and improve your experience. You can choose which cookies to allow. [Privacy Policy](#)

Accept All

Decline

TECHSY		
Clean, maintainable pipeline architecture	Haystack + Qdrant	Optimized framework + performant open-source DB
Measuring retrieval quality	RAGAS + Langfuse	Batch evaluation + real-time monitoring, both free
Maximum flexibility for complex chains	LangChain + Pinecone + Cohere Rerank	Most integrations + zero-ops DB + precision reranking

For a broader view of AI infrastructure choices beyond RAG, see [our AI stack guide for SaaS](#).

Need Something Custom?

Off-the-shelf RAG stacks work for most use cases, but some projects need more. Multi-modal retrieval across images and text. Domain-specific reranking models fine-tuned on your data. Hybrid architectures that combine RAG with agentic reasoning.

At Techsy, we've built custom RAG pipelines for clients who outgrew their starter stacks and needed something tailored, from selecting the right embedding model for niche domains to designing evaluation frameworks that measure what actually matters for their users. If you're also evaluating the LLM routing layer that sits in front of your RAG pipeline, our [LLM gateway tools comparison](#) covers Portkey, LiteLLM, and OpenRouter side by side.

If you're past the prototype stage and need help optimizing or building a production RAG system, [reach out for a free consultation](#). We'll start with your data, your queries, and your accuracy requirements, not a generic template. [See our AI integration services](#).

FAQ



We use cookies to analyze site traffic and improve your experience. You can choose which cookies to allow. [Privacy Policy](#)

Accept All

Decline

TECHSY

The best RAG tools depend on your pipeline stage. LangChain for orchestration, Pinecone or Weaviate for vector storage, OpenAI text-embedding-3-large for embeddings, Cohere Rerank for precision, and RAGAS for evaluation. See the rankings table above for our full ranked list with alternatives.

Should I use LangChain or LlamaIndex for RAG?

LangChain if you need maximum flexibility and a huge ecosystem for complex multi-step pipelines. LlamaIndex if your RAG is 80%+ document retrieval, its 300+ data connectors and purpose-built query engines make document-heavy use cases significantly easier. Many teams actually use both together.

What is the best vector database for RAG?

Weaviate for production RAG that needs hybrid search (vector + keyword). Pinecone if you want zero-ops managed infrastructure. Chroma for prototyping. The "best" depends entirely on whether you prioritize features, simplicity, or operational overhead.

Is Pinecone worth paying for over open-source alternatives?

Yes, if your team lacks DevOps capacity and needs enterprise compliance features. No, if you have the infrastructure skills to run Weaviate or Qdrant yourself, you'll get comparable performance at a fraction of the cost. Model your expected query volume and storage needs before deciding.

Do I need a reranking model for RAG?

Not for prototypes or small document sets. For production RAG where answer accuracy matters, reranking typically improves relevancy by 10-20%. Cohere Rerank is the easiest to add (one API call), or self-host Jina Reranker v2 to avoid API costs.



We use cookies to analyze site traffic and improve your experience. You can choose which cookies to allow. [Privacy Policy](#)

Accept All

Decline

Use RAGAS to measure four key metrics: context precision (are the retrieved docs

TECHSY

answer stick to the sources?), and answer relevancy (does it actually answer the

question?). Run evaluations on real queries from your domain, not synthetic benchmarks.

What's the cheapest way to build a RAG pipeline?

LangChain + Chroma + OpenAI text-embedding-3-small + RAGAS. Total cost under \$10/month, mostly embedding API calls. Swap OpenAI embeddings for self-hosted BGE and the entire stack is free except compute.

Can I use Haystack instead of LangChain?

Absolutely. Haystack's pipeline architecture is arguably cleaner and more maintainable. The trade-off is ecosystem size, fewer tutorials, integrations, and community resources. If you value explicit, well-structured code over having a pre-built integration for everything, Haystack is the better choice.

How do Weaviate and Pinecone compare for RAG?

Weaviate offers native hybrid search and is open-source (self-host or use managed cloud). Pinecone is fully managed with zero operational overhead but proprietary. Weaviate is typically cheaper, more feature-rich for RAG-specific use cases, and avoids vendor lock-in. Pinecone wins on simplicity and enterprise compliance.

What RAG monitoring tools should I use in production?

Langfuse for open-source, framework-agnostic monitoring with tracing, prompt management, and evaluation. LangSmith if you're already deep in the LangChain ecosystem and want tightly integrated tracing. Both offer free tiers. Pair either one with RAGAS for periodic batch evaluations of retrieval quality.



We use cookies to analyze site traffic and improve your experience. You can choose which cookies to allow. [Privacy Policy](#)

Accept All

Decline

unmatched ecosystem, LangGraph for agentic pipelines, and LangSmith for end-to-end

TECHSY

retrieval, its 500+ data connectors and structured query engines are purpose-built for

that job. If your team values clean, auditable architecture over ecosystem breadth, Haystack v2 is worth a hard look. The "best" framework depends on your retrieval complexity, team size, and whether you need agent orchestration alongside basic RAG.

Which RAG tool should I use?

Start with your bottleneck. Prototyping? Use Chroma + LangChain, zero config, running in 15 minutes. Need production vector search? Weaviate if you want open-source hybrid search; Pinecone if you want fully managed infrastructure and can pay the premium. Answer quality too low after retrieval? Add Cohere Rerank before you change anything else, it's one API call and typically adds 10-20% relevancy. Not sure if your pipeline is actually working? RAGAS gives you four metrics that matter. Most production teams end up with three to four of these tools in combination, not just one.

Sources

[LangChain Documentation](#)

[LlamaIndex Documentation](#)

[Pinecone Pricing](#)

[Weaviate Documentation](#)

[Chroma Documentation](#)

[Cohere Rerank Documentation](#)

[RAGAS Documentation](#)

[Haystack Documentation](#)

[OpenAI Embeddings Guide](#)

[Langfuse Documentation](#)



We use cookies to analyze site traffic and improve your experience. You can choose which cookies to allow. [Privacy Policy](#)

Accept All

Decline

best rag tools

rag frameworks

vector database

langchain

llamaindex

pinecone

TECHSY

Share this article

 Copy link

 in

 X

 ✉



We use cookies to analyze site traffic and improve your experience. You can choose which cookies to allow. [Privacy Policy](#)

Accept All

Decline

TECHSY



We use cookies to analyze site traffic and improve your experience. You can choose which cookies to allow. [Privacy Policy](#)

Accept All

Decline

TECHSY



We use cookies to analyze site traffic and improve your experience. You can choose which cookies to allow. [Privacy Policy](#)

Accept All

Decline

TECHSY

RESOURCES

[SEE ALL](#)

The Software Procurement Playbook

A repeatable framework for buying software without burning six months and a million dollars on the wrong platform.

The Architecture Decision Playbook

A practical framework for picking your stack: when to build vs. buy, monolith vs. microservices, and how to avoid resume-driven design.

The Vendor Selection Playbook

How to pick the right development partner (agency, freelancer, in-house) without overpaying or shipping a half-built product.

CLAUDE SKILLS

[SEE ALL](#)

New Post

Full SEO blog pipeline: research, brief, write, validate, image, translate, publish to Sanity. Autonomous from start to finish.

Content Refresh

Audit a stale post, find decay drivers, and ship a SERP-aligned refresh without losing existing rankings.

SEO Audit

Site-wide SEO audit with prioritized fix list: technical, on-page, and EEAT signals.

AI AUTOMATIONS

[SEE ALL](#)

Security Auditor

Weekly SCA + IaC scan with prioritized fix PRs.

Cold Email Writer

Generates first-touch emails grounded in one specific public detail.

Lead Research Agent

Enrich an email into a profile, score fit, alert in Slack.



We use cookies to analyze site traffic and improve your experience. You can choose which cookies to allow. [Privacy Policy](#)

[Accept All](#)[Decline](#)

TECHSY

Enterprise Solutions

Mobile Apps

Web Applications

CRM Systems

AI Integration

ERP Solutions

Voice Agents

Process Automation

Cybersecurity

LIBRARY

Resources

Blog

Portfolio

COMMUNITY

AI Automations

Claude Skills

TOOLS

Mobile App Cost Calculator

OpenAI / LLM API Cost Calculator

MVP Cost Calculator

Voice AI Agent Cost Calculator

COMPANY

About

Partners

Contact

LEGAL

Privacy Policy

Terms of Service

Cookie Policy

TECHSY

© 2026 Techsy. All rights reserved.



We use cookies to analyze site traffic and improve your experience. You can choose which cookies to allow. [Privacy Policy](#)

Accept All

Decline